Fireworks AI

Meta Llama / Llama 3.2 3B Instruct

accounts/fireworks/models/llama-v3p2-3b-instruct

Serverless LLM Tunable Chat

Open in Playground

Details Playground

Llama 3.2 3B instruct is a lightweight, multilingual model from Meta. The model is designed for efficiency and offers substantial latency and cost improvements compared to larger models. Example use cases for the model include query and prompt rewriting and writing assistance

Serverless API

Llama 3.2 3B Instruct is available via Fireworks' serverless API, where you pay per token. There are several ways to call the Fireworks API, including Fireworks' Python client, the REST API, or OpenAI's Python client.

See below for easy generation of calls and a description of the raw REST API for making API requests. See the Querying text models docs for details.

Try it

API Examples

Generate a model response using the chat endpoint of **llama-v3p2-3b-instruct.**
API reference

Python Typescript Java Go Shell Chat Completion

```
import requests
import json

url = "https://api.fireworks.ai/inference/v1/chat/completions"
payload = {
  "model": "accounts/fireworks/models/llama-v3p2-3b-instruct",
  "max_tokens": 16384,
  "top_p": 1,
  "top_k": 40,
  "presence_penalty": 0,
  "frequency_penalty": 0,
  "temperature": 0.6,
  "messages": [
    {
      "role": "user",
      "content": "Hello, how are you?"
    }
  ]
}
headers = {
  "Accept": "application/json",
  "Content-Type": "application/json",
  "Authorization": "Bearer <API_KEY>"
}
requests.request("POST", url, headers=headers, data=json.dumps(payload))
```

Fine-tuning

Llama 3.2 3B Instruct can be fine-tuned on your data to create a model with better response quality. Fireworks uses low-rank adaptation (LoRA) to train a model that can be served efficiently at inference time.

See the Fine-tuning guide for details.

Fine-tune this model

On-demand deployments

On-demand deployments allow you to use **Llama 3.2 3B Instruct** on dedicated GPUs with Fireworks' high-performance serving stack with high reliability and no rate limits.

See the On-demand deployments guide for details.

Deploy this Base model

Open in model playground

Finetune this model

Deploy this model

Model Details

Created by	dzhulgakov@fireworks.ai
Created	9/18/2024
Visibility	Public
Kind	Base model
Model size	3B parameters
Fine-tuning	Supported
Serverless LoRA Deployment	Supported
Function Calling	Not supported
Provider	Meta Llama
Hugging Face	Visit link